

# Project: TMDb Movies project UDACITY

By Ifalore Simeon

## Table of Contents

- [Introduction](#)
- [Data Wrangling](#)
- [Data Cleaning](#)
- [Exploratory Data Analysis](#)
- [Conclusions](#)

## Introduction

This data set contains information about 10,000 movies collected from The Movie Database (TMDb), including user ratings and revenue. Certain columns, like 'cast' and 'genres', contain multiple values separated by pipe (|) characters. There are some odd characters in the 'cast' column. Don't worry about cleaning them. You can leave them as is. The final two columns ending with "\_adj" show the budget and revenue of the associated movie in terms of 2010 dollars, accounting for inflation over time.


The final two columns ending with "\_adj" show the budget and revenue of the associated movie in terms of 2010 dollars, accounting for inflation over time.

### SOURCE TO DATASET

[https://d17h27t6h515a5.cloudfront.net/topher/2017/October/59dd1c4c\\_tmdb-movies/tmdb-movies.csv](https://d17h27t6h515a5.cloudfront.net/topher/2017/October/59dd1c4c_tmdb-movies/tmdb-movies.csv) ([https://d17h27t6h515a5.cloudfront.net/topher/2017/October/59dd1c4c\\_tmdb-movies/tmdb-movies.csv](https://d17h27t6h515a5.cloudfront.net/topher/2017/October/59dd1c4c_tmdb-movies/tmdb-movies.csv)).

### Questions:

1. Top 5 most expensive movies
2. What is the average runtime, budget, revenue of the movies
3. What is the movie with the highest and lowest profit
4. Which year had most profitable movies
5. The most popular movie before the movie with the highest profit was released
6. Is there a correlation between how popular a movie is and the profit?

```
In [1]:  # Use this cell to set up import statements for all of the packages that you  
# plan to use.  
  
# Remember to include a 'magic word' so that your visualizations are plotted  
# inline with the notebook. See this page for more:  
# http://ipython.readthedocs.io/en/stable/interactive/magics.html  
  
import pandas as pd  
import numpy as np  
from datetime import datetime  
import matplotlib.pyplot as plt  
%matplotlib inline
```

## Data Wrangling

The data wrangling process will be done using pandas. Some methods will be used to explore the dataset and draw intuition about the dataset to understand the shape, datatype and other parameters\*\* First of all I will load my data and use several codes to draw insight as to the shape, datatypes, NAN values etc present in the data

In [2]:

```
df = pd.read_csv('tmdb-movies.csv')
# Looking at null values and datatypes so I can know what cleanin techniques
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     10866 non-null  int64
1   imdb_id               10856 non-null  object
2   popularity            10866 non-null  float64
3   budget                10866 non-null  int64
4   revenue               10866 non-null  int64
5   original_title        10866 non-null  object
6   cast                  10790 non-null  object
7   homepage              2936 non-null   object
8   director              10822 non-null  object
9   tagline               8042 non-null   object
10  keywords              9373 non-null   object
11  overview              10862 non-null  object
12  runtime               10866 non-null  int64
13  genres                10843 non-null  object
14  production_companies  9836 non-null   object
15  release_date          10866 non-null  object
16  vote_count            10866 non-null  int64
17  vote_average          10866 non-null  float64
18  release_year          10866 non-null  int64
19  budget_adj            10866 non-null  float64
20  revenue_adj           10866 non-null  float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB
```

In [3]:

```
#Exploring the shape of the dataset
df.shape
```

Out[3]: (10866, 21)

## observations

- The data set has "10866" rows and "21" columns
- There are some null values in the data set
- data type looks consistent but release\_date is stored as a string
- budget and revenue are floats
- there are some 0 values in our data sets

In [4]:

```
# Confirming that release_date is stored as a string
type(df['release_date'][0])
```

Out[4]: str

In [5]: `# exploring the first 5 rows to have a sense of my data`  
`df.head()`

Out[5]:

	id	imdb_id	popularity	budget	revenue	original_title	cast	
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	
1	76341	tt1392190	28.419936	150000000	378436354	Mad Max: Fury Road	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	
2	262500	tt2908446	13.112507	110000000	295238201	Insurgent	Shailene Woodley Theo James Kate Winslet Ansel...	http://m
3	140607	tt2488496	11.173104	200000000	2068178225	Star Wars: The Force Awakens	Harrison Ford Mark Hamill Carrie Fisher Adam D...	
4	168259	tt2820852	9.335014	190000000	1506249360	Furious 7	Vin Diesel Paul Walker Jason Statham Michelle ...	

5 rows × 21 columns

In [6]: `# checking for duplicates in the data`  
`sum(df.duplicated())`

Out[6]: 1

In [7]: `df[df.duplicated()]`

Out[7]:

	id	imdb_id	popularity	budget	revenue	original_title	cast	homepage
2090	42194	tt0411951	0.59643	30000000	967000	TEKKEN	Jon Foo Kelly Overton Cary-Hiroynuki Tagawa Ian...	NaN

1 rows × 21 columns

## Data Cleaning -

This list below shows the different cleaning techniques that will be employed to make the dataset appropriate for Exploratory data Analysis

- Changing release date column from string to datetime data type.
- Removing the duplicated values.
- Changing format of budget and revenue columns.
- Remove unnecessary columns such as 'imdb\_id', 'budget\_adj', 'revenue\_adj', 'homepage', 'production\_company', 'keywords', 'overview', 'production\_companies', 'vote\_count' and 'vote\_average'.
- replace all the values from '0' to NAN in 'budget' and 'revenue' columns, then removing them.

### Changing release date column from string to date type data.

To change the release\_date to datetime I will have to use the datetime function of pandas

```
In [8]: df['release_date'] = pd.to_datetime(df['release_date'])  
# checking to see that the datatypes has changed to timestamp for 'release_date'  
type(df['release_date'][0])
```

```
Out[8]: pandas._libs.tslibs.timestamps.Timestamp
```

Now I need to confirm that my release\_date is now a date datatype

In [9]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    10866 non-null  int64
1   imdb_id                             10856 non-null  object
2   popularity                           10866 non-null  float64
3   budget                              10866 non-null  int64
4   revenue                             10866 non-null  int64
5   original_title                       10866 non-null  object
6   cast                                10790 non-null  object
7   homepage                             2936 non-null  object
8   director                             10822 non-null  object
9   tagline                              8042 non-null  object
10  keywords                             9373 non-null  object
11  overview                             10862 non-null  object
12  runtime                              10866 non-null  int64
13  genres                               10843 non-null  object
14  production_companies                 9836 non-null  object
15  release_date                         10866 non-null  datetime64[ns]
16  vote_count                           10866 non-null  int64
17  vote_average                         10866 non-null  float64
18  release_year                         10866 non-null  int64
19  budget_adj                           10866 non-null  float64
20  revenue_adj                          10866 non-null  float64
dtypes: datetime64[ns](1), float64(4), int64(6), object(10)
memory usage: 1.7+ MB
```

## Removing duplicate data

Before proceeding with the analysis some duplicate values need to be dropped so they don't affect the outcome of the analysis

In [10]: `# removing duplicate data`  
`df.drop_duplicates(keep='first', inplace=True)`  
`# checking for duplicated value`  
`df.duplicated().sum()`

Out[10]: 0

```
In [11]: df.nunique()
```

```
Out[11]: id                10865  
         imdb_id          10855  
         popularity       10814  
         budget           557  
         revenue          4702  
         original_title   10571  
         cast             10719  
         homepage         2896  
         director         5067  
         tagline           7997  
         keywords          8804  
         overview         10847  
         runtime           247  
         genres            2039  
         production_companies 7445  
         release_date       5909  
         vote_count        1289  
         vote_average        72  
         release_year        56  
         budget_adj         2614  
         revenue_adj        4840  
         dtype: int64
```

**Changing format of budget and revenue columns.**

```
In [12]: ▶ change_type=['budget', 'revenue']

#changing data type
df[change_type]=df[change_type].applymap(np.int64)

#printing data types of the dataset to see the changed information
df.dtypes
```

```
Out[12]: id                int64
imdb_id                object
popularity            float64
budget                int64
revenue              int64
original_title        object
cast                 object
homepage             object
director             object
tagline             object
keywords             object
overview             object
runtime              int64
genres               object
production_companies  object
release_date          datetime64[ns]
vote_count            int64
vote_average          float64
release_year          int64
budget_adj            float64
revenue_adj           float64
dtype: object
```

## Removing unnecessary columns

Since I won't be using all columns in my analysis, I will be dropping the columns that I won't be needing at all through the course of the EDA because they are not relevant to answering the posed questions



```
In [13]: ▶ #List of columns I want to remove
col_to_del = ['imdb_id', 'budget_adj', 'revenue_adj', 'homepage', 'keywords',
#deleting the columns from the database
df = df.drop(col_to_del, 1)

# previewing the new dataset
df.head(3)
```

C:\Users\Simeon\AppData\Local\Temp\ipykernel\_2248\2768113811.py:4: FutureWarning: In a future version of pandas all arguments of DataFrame.drop except for the argument 'labels' will be keyword-only.

```
df = df.drop(col_to_del, 1)
```

Out[13]:

	id	popularity	budget	revenue	original_title	cast	director	tagline
0	135397	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	Colin Trevorrow	The park is open
1	76341	28.419936	150000000	378436354	Mad Max: Fury Road	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	George Miller	What a Lovely Day
2	262500	13.112507	110000000	295238201	Insurgent	Shailene Woodley Theo James Kate Winslet Ansel...	Robert Schwentke	One Choice Car Destroy You

## Replacing the '0' in our column to NAN so we can drop it

After exploring the first and last few rows in budget and revenue columns, there are several entries have '0' value, so firstly I will convert those values to NAN then I will remove them from the dataset.

```
In [14]: ▶ # creating a list of revenue and budget columns
temp_list=['budget', 'revenue']

#this will replace all the value from '0' to NAN in the list
df[temp_list] = df[temp_list].replace(0, np.NAN)

#Removing all the row which has NaN value in temp_list
df.dropna(subset = temp_list, inplace = True)
```

After removing them we need to check to see that the amount of columns have reduced to be sure the affected columns have been dropped

In [15]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3854 entries, 0 to 10848
Data columns (total 12 columns):
#   Column              Non-Null Count  Dtype
---  -
0   id                  3854 non-null   int64
1   popularity          3854 non-null   float64
2   budget              3854 non-null   float64
3   revenue             3854 non-null   float64
4   original_title      3854 non-null   object
5   cast                3850 non-null   object
6   director            3853 non-null   object
7   tagline             3574 non-null   object
8   runtime             3854 non-null   int64
9   genres              3854 non-null   object
10  release_date        3854 non-null   datetime64[ns]
11  release_year        3854 non-null   int64
dtypes: datetime64[ns](1), float64(3), int64(3), object(5)
memory usage: 391.4+ KB
```

## Exploratory Data Analysis

### Research Question 1 (Top 5 most expensive movies)

To get the top 5 most expensive movies we need to sort the dataset based on budget in descending order. This will give us the list from most expensive to the least expensive

In [16]:

```
expensive_movies = df.sort_values(by='budget',ascending=False).head()
expensive_movies
```

Out[16]:

	id	popularity	budget	revenue	original_title	cast	director
<b>2244</b>	46528	0.250540	425000000.0	1.108757e+07	The Warrior's Way	Kate Bosworth Jang Dong-gun Geoffrey Rush Dann...	Sngmoo Lee
<b>3375</b>	1865	4.955130	380000000.0	1.021683e+09	Pirates of the Caribbean: On Stranger Tides	Johnny Depp PenÃ©lope Cruz Geoffrey Rush Ian M...	Rob Marshall
<b>7387</b>	285	4.965391	300000000.0	9.610000e+08	Pirates of the Caribbean: At World's End	Johnny Depp Orlando Bloom Keira Knightley Geof...	Gore Verbinski
<b>14</b>	99861	5.944927	280000000.0	1.405036e+09	Avengers: Age of Ultron	Robert Downey Jr. Chris Hemsworth Mark Ruffalo...	Joss Whedon
<b>6570</b>	1452	1.957331	270000000.0	3.910812e+08	Superman Returns	Brandon Routh Kevin Spacey Kate Bosworth James...	Bryan Singer



```
In [17]: ▶ least_expensive = df.sort_values(by='budget')
least_expensive.head()
```

Out[17]:

	id	popularity	budget	revenue	original_title	cast	director	tagline
3581	59296	0.520430	1.0	1378.0	Love, Wedding, Marriage	Mandy Moore Kellan Lutz Jessica Szohr Autumn F...	Dermot Mulroney	Here comes the ride.
2618	39964	0.090186	1.0	100.0	Lost & Found	David Spade Sophie Marceau Ever Carradine Step...	Jeff Pollack	A comedy about a guy who would do anything to ...
8944	14373	0.464188	2.0	16.0	Death Wish 2	Charles Bronson Jill Ireland Vincent Gardenia ...	Michael Winner	First His Wife. Now His Daughter. It's Time To...
10050	20701	0.317091	3.0	16.0	Tales from the Darkside: The Movie	Rae Dawn Chong Christian Slater Deborah Harry ...	John Harrison	From the depths of four twisted minds.
2398	39356	0.028456	3.0	43.0	Boy	James Rolleston Craig Hall Taika Waititi Te Ah...	Taika Waititi	Summer, Girls, Gangs, Drugs ... its not easy b...



Movie ID 2244 The Warrior's Way shows the largest budget i.e 425000000

Whereas Love, Wedding, Marriage with id no. 3581 shows the smallest budget i.e 1 dollar

## Research Question 2 (What is the average runtime, budget, revenue of the movies)

To avoid repetition I am going to create a simple function called avg i.e average to calculate the average runtime, budget and revenue. This will give us an idea of where majority of the movies fall in terms of budget, revenue and runtime.

```
In [18]: # defining a function to find average of a column  
def avg(column):  
    return df[column].mean()
```

```
In [19]: # calculating the average runtime  
avg(['runtime'])
```

```
Out[19]: runtime    109.220291  
dtype: float64
```

No unit was given for runtime, budget and revenue so we would assume runtime is in minutes, revenue and budget are in dollars

```
In [20]: # calculating the average budget. in this case a print function was added to  
print(round(avg(['budget'])));
```

```
budget    37203697.0  
dtype: float64
```

```
In [21]: # calculating the average revenue  
print(round(avg(['revenue'])));
```

```
revenue    107686616.0  
dtype: float64
```

On the average the runtime, budget and revenue are 109.2 mins, 37203697 dollars, 107686616 dollars

### Research Question 3 (What is the movie with the highest and lowest profit)

We don't have a column represents the total profits of the movies in our dataset to find the most and least profit movies so I will answer this question in two steps:

#### A. Calculating the profit of each movie

I'm going to add a new column to the dataset represents the total profit every movie in our dataset made by subtracting the revenue from the budget.

```
In [22]: # Finding profit  
df['profit'] = df['revenue']-df['budget']
```

#### B. From the profit column calculate movies with most and least profit

In [23]:

```
# Which movie has highest profit?
highest_profit_movie = df.sort_values(by='profit',ascending=False)
highest_profit_movie.iloc[0]
```

```
Out[23]: id                19995
popularity              9.432768
budget                237000000.0
revenue              2781505847.0
original_title          Avatar
cast      Sam Worthington|Zoe Saldana|Sigourney Weaver|S...
director              James Cameron
tagline              Enter the World of Pandora.
runtime                162
genres      Action|Adventure|Fantasy|Science Fiction
release_date      2009-12-10 00:00:00
release_year              2009
profit              2544505847.0
Name: 1386, dtype: object
```

```
In [24]: lowest_profit_movie = df.sort_values(by='profit')
lowest_profit_movie.iloc[0]
```

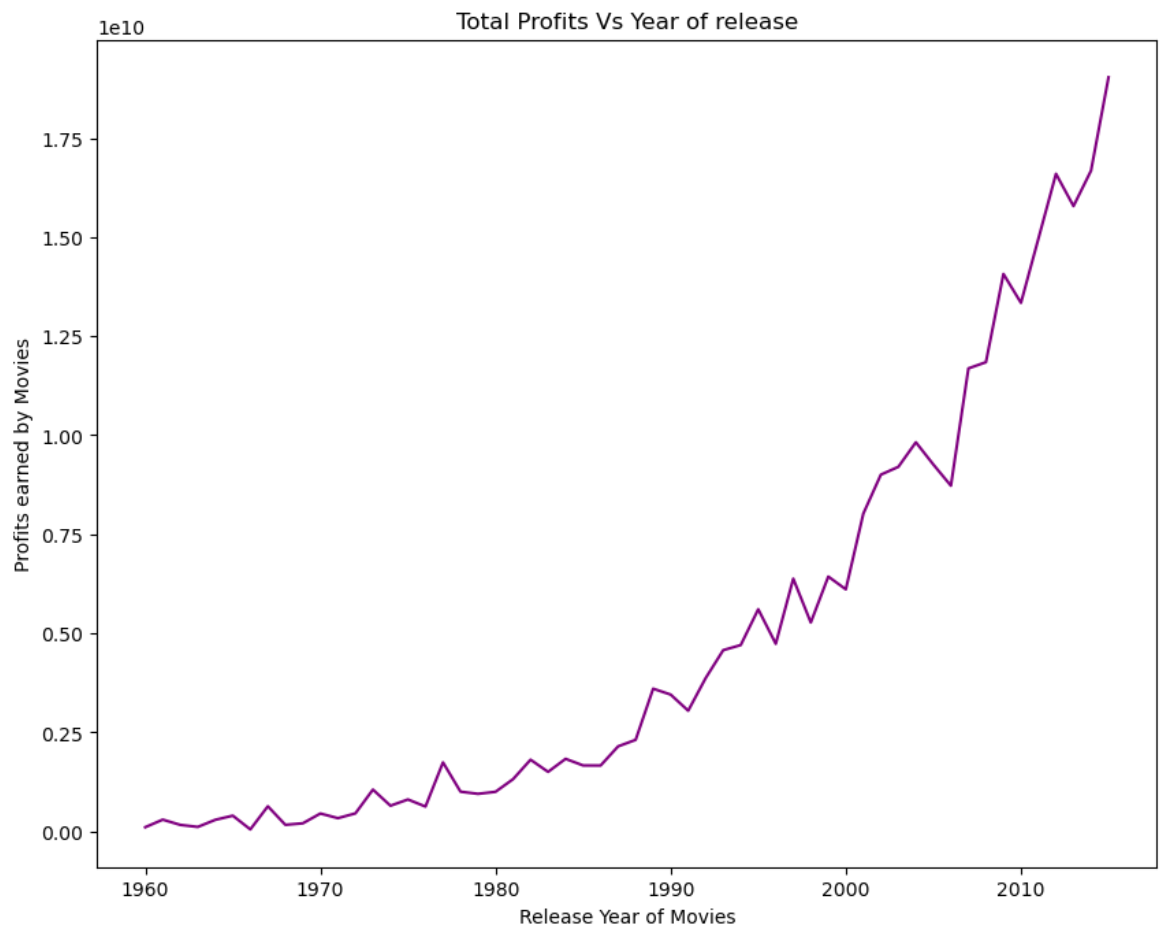
```
Out[24]: id                46528
popularity              0.25054
budget                425000000.0
revenue              11087569.0
original_title      The Warrior's Way
cast      Kate Bosworth|Jang Dong-gun|Geoffrey Rush|Dann...
director              Sngmoo Lee
tagline      Assassin. Hero. Legend.
runtime                100
genres      Adventure|Fantasy|Action|Western|Thriller
release_date      2010-12-02 00:00:00
release_year              2010
profit              -413912431.0
Name: 2244, dtype: object
```

Movie Avatar with id 19995 shows the highest earned profit i.e 237000000.

WhereasThe Warrior's Way movie with id 46528 shows the lowest earned profit i.e -413912431

## Research Question 4 (Which year had most profitable movies)

```
In [25]: #We will be using Line plot for this analysis  
#Since we want to know the profits of movies for every year therefore we have  
  
profits_year = df.groupby('release_year')['profit'].sum()  
  
#figure size(width, height)  
plt.figure(figsize=(10,8), dpi = 100)  
  
#Labelling x-axis  
plt.xlabel('Release Year of Movies')  
#Labelling y-axis  
plt.ylabel('Profits earned by Movies')  
#title of the plot  
plt.title('Total Profits Vs Year of release')  
  
#plotting the graph  
plt.plot(profits_year, color = 'purple')  
plt.show()
```



```
In [26]: ▶ #To find which year made the highest profit?
profits_year.idxmax()
```

Out[26]: 2015

Graphically we cannot easily spot that 2015 is the year with the highest profit but using the `idxmax()` helps confirm that. However an important observation from the graph is that we can see an upwards trend of profit Year on Year

## Research Question 5 (The most popular movie before the movie with the highest profit was released)

Avatar was the movies with the highest profit in 2009 but before then what movie was the most popular?

```
In [27]: ▶ # in this case the highest profit earning movie was Avatar released in 2009
popular_movie_2010_idx = df[df["release_year"]<=2009]["popularity"].idxmax()
df.loc[popular_movie_2010_idx]
```

```
Out[27]: id                                     11
popularity                                     12.037933
budget                                       11000000.0
revenue                                   775398007.0
original_title                               Star Wars
cast      Mark Hamill|Harrison Ford|Carrie Fisher|Peter ...
director                                George Lucas
tagline      A long time ago in a galaxy far, far away...
runtime                                           121
genres      Adventure|Action|Science Fiction
release_date      1977-03-20 00:00:00
release_year      1977
profit      764398007.0
Name: 1329, dtype: object
```

The most popular movie was Star Wars released in 1977 before Avatar

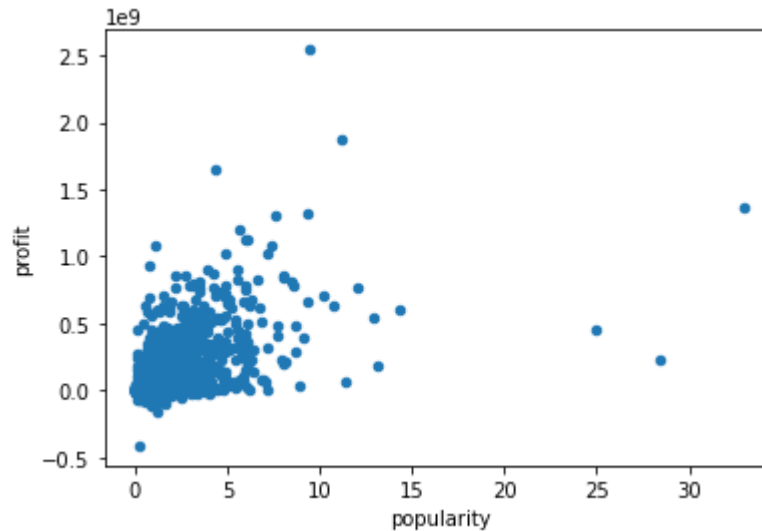
## Research Question 5 (Is there a correlation between popularity and profit?)

We have looked at Avatar and Star Wars in previous sections but does the popularity of a movie affect how much profit is earned from it?

Firstly, I will create a scatter plot between popularity and profit to visually see the correlation and then use the `corr()` to get the exact value



```
In [28]: # creating a scatterplot between popularity and profit
df.plot(x = 'popularity', y = 'profit', kind = 'scatter');
```



```
In [29]: # getting the quantitative value for the correlation between profit and popul
corr = df['popularity'].corr(df['profit'])
corr
```

Out[29]: 0.5962013674920571

## Conclusions

To summarize the whole analysis:

- **SUMMARY**

The top 5 most expensive movies are:

- The Warrior's Way
- Pirates of the Caribbean: On Stranger Tides
- Pirates of the Caribbean: At World's End
- Avengers: Age of Ultron
- Superman Returns

The average:

- runtime is 109.220291
- budget is 37203697.0
- revenue is 107686616.0

- Movie with the highest profit is 'Avatar', While Movie with the lowest profit is 'The Warrior's Way'

- The year with the most profit is 2015 While there is a 0.59 correlation between popularity and profit

### **CONCLUSION**

- We can observe an ongoing upwards trend in Total profit of movies every year.
- There is a moderate relationship between popularity of a movie and profit made
- The most expensive movies made belong to the genres of Action|Adventure|Fantasy/Sci-Fi

### **LIMITATIONS**

- Since the budget and revenue column do not have currency unit, it might be possible different movies have budget in different currency according to the country they were produced in. So an inconsistency appears here which can state the complete analysis wrong. Dropping the rows with missing values might have also affected the overall analysis.